Robust Performance Control for Web Applications in the Cloud

Corina Stratan, Guillaume Pierre, Hector Fernandez, Eliana Tirsa and Valentin Cristea

Vrije Universiteit Amsterdam and University Politehnica of Bucharest

ONE OF the major innovations provided by Cloud computing platforms is their pay-per-use model where applications can request and release resources at any time according to their needs — and pay only for the resources they actually used. This business model is particularly favorable for application domains where workloads vary widely over time, such as the domain of Web applications.

However, provisioning the right quantities of resources for a Web application is not a simple task. Web applications are usually composed of multiple types of components such as load balancers, Web servers, application servers and database servers. Complex performance behavior of these components make it difficult to find the optimal resource allocation, even when the application workload is perfectly stable. The magnitude of the problem is further increased by the fact that Web application workloads are often very unstable and hard to predict. In the case of a sudden load increase there is a necessary tradeoff between reacting as early as possible to minimize the duration when the application underperforms because of insufficient processing capacity, and a slower approach to avoid situations where the load has already decreased when the new resources become available.

Faced with this difficult scientific challenge, the academic community has proposed a wide range of sophisticated resource provisioning algorithms [?, ?, ?]. However, we observe a discrepancy between these academic propositions and the simple mechanisms that are currently available to cloud customers. These mechanisms are usually based on defining lower or upper thresholds on resource utiliza-tion, and reaching one of the thresholds will trigger a resource provisioning action.

We postulate three possible reasons why sophisticated techniques have not been more widely deployed: *(i)* the gains of using sophisticated provisioning strategies are too low to be worth the effort; *(ii)* implementing these techniques is a difficult exercise, which is why real cloud systems rely on simpler techniques; and *(iii)* academic approaches mostly focus on unrealistic evaluations using simple applications and artificial workloads [4, 8, 19].

CORINA: *I reformulated this because I think it is better to be more constructive and state that we not only look for the cause, but we also try to see if there is a solution; also, it could be that more than one of these 3 possible causes are real* This paper aims to investigate which of these possible causes are the real problems, and to find whether it is possible to implement an automatic scaling algorithm which provides better results than the simple ones without being too complex. We do this by implementing a provisioning system in realistic conditions, and reporting on *(i)* how hard implementation was; and *(ii)* potential gains from using this technique compared to the simple solution. To achieve this, we designed and implemented a resource provisioning technique in ConPaaS, an open source platform-as-a-service environment for hosting cloud applications [14]. Our technique uses several levels of thresholds to predict future performance degradations, workload trend detection to better handle traffic spikes and dynamic load balancing weights to handle resources heterogeneity. To test our system in realistic settings, we deployed a copy of Wikipedia

and replayed a fraction of the real access traces to its official web site.

ConPaaS overview

ConPaaS is an open-source runtime environment for hosting applications in Cloud infrastructures [14]. Within the Cloud computing paradigm, ConPaaS belongs to the platform-as-a-service family, in which a variety of systems aim to simplify the deployment of applications in the Cloud. Using ConPaaS, developers can now focus their attention on application-specific concerns rather than making their applications suitable for the cloud.

Architecture

In ConPaaS, an application is designed as a composition of one or more elastic and distributed components, called *services*. Each service is dedicated to host a particular type of functionality of an application. At the moment, ConPaaS supports six different types of services: two web application hosting services respectively specialized for hosting PHP and JSP applications; a MySQL database service; a NoSQL database service built around the Scalarix key-value store; a MapReduce service; and a Task-Farming service for high-performance batch processing.

The architecture of ConPaaS services comprises two main building blocks: agents and managers.

- **Agent**: The agent VMs (virtual machines) host the needed components to provide the service-specific functionality. Based on the performance requirements or the application workload, the number of agent VMs hosting these components can grow or shrink on demand. Thus, as an example, the minimal configuration of the PHP web hosting service consists of a single agent VM. If necessary, the service can progressively grow to multiple agent VMs hosting its components: load-balancer, web servers for static pages and/or PHP servers for dynamic pages.

- **Manager**: For each service, there is only one manager VM. The manager is in charge of centralizing monitoring data and controlling the allocation of resources assigned to the service.

Figure 1 shows the architecture of the ConPaaS deployment for a typical PHP web application backed by a MySQL database. The application is deployed using two ConPaaS services: the PHP service and the MySQL service.

Hosting Elastic Applications

The main features that distinguish ConPaaS from other PaaS systems are its approach for autonomous application scaling and its interoperability with several private and public IaaS clouds. To provide such autonomous scaling capabilities, ConPaaS includes a monitoring data analysis mechanism and a resource provisioning system.

For monitoring in ConPaas we used Ganglia [6], a scalable distributed monitoring system. The monitoring data is collected by the manager VM, which runs for this purpose a Ganglia meta daemon (gmetad). We implemented modules that extend Ganglia's standard set of monitoring metrics, by adding a number of service-specific metrics. For the PHP web hosting service, we added monitoring metrics for measuring the request rate and response time both at the static web server and at the PHP server.

Our approach for automatic resource provisioning is to implement customized scaling policies for each type of service, taking into account the particularities of the service (e.g. the structure of the service deployment, the complexity of the requests etc.) and also the service-specific monitoring data. So far our focus has been on the web hosting service, for which we implemented a number of automatic resource provisioning mechanisms discussed in Section .

The Wikipedia application

To evaluate the resource provisioning algorithms implemented in ConPaaS, we prepared a realistic testing scenario: we used a copy of Wikipedia as the
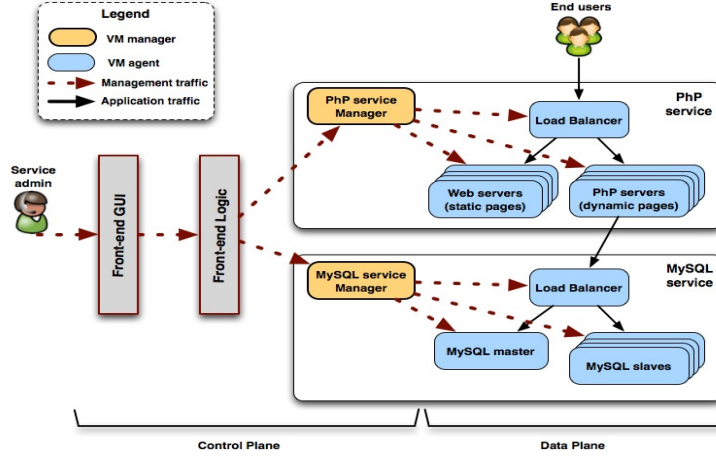
Figure 1: ConPaaS system architecture.

web application, and replayed a fraction of the access traces to Wikipedia's official web site.

Wikipedia uses MediaWiki [20], a PHP-based wiki software package. To deploy the Wikipedia services on ConPaaS, we composed two different services: the PHP web hosting and MySQL service. In the MySQL service, we installed a complete copy of the English Wikipedia database from 2008, which contains about 30GB in Wikipedia articles. In the PHP service, the configuration was composed of one load balancer, one static web server and one or more PHP servers. ConPaaS uses nginx as static web server and also as load balancer (the HttpProxy module of nginx). The PHP requests are handled through PHP-FPM (PHP FastCGI Process Manager).

In order to benchmark ConPaaS when hosting the Wikipedia services, we used Wikibench [5], a tool developed by our group for benchmarking Wikipedia servers. Wikibench replays Wikipedia access traces or smaller samples of the traces, providing various sampling mechanisms. We used anonymized Wikipedia access traces, which are published by the WikiMedia foundation, and for our experiments we replayed a sample of 10% from the traces. The traces contain two types of requests: requests for static pages or files, and dynamic PHP requests. The static requests are executed much faster in comparison with the dynamic requests, so in this paper we will discuss mostly the dynamic requests, which pose more performance problems.

CORINA: *I wouldn't say the traces are heterogeneous in terms of workload mix. Usually workload mix refers to different types of requests – e.g. log in, read article, update article etc. In our case it's most of the time the same type of request (read an article) but with articles of various complexity.*

As an example, in Figure 2, we show the PHP workload sampled from one trace, as the number of PHP requests per minute during approximately one day. One aspect that we noticed about the workload is that the requests are very heterogeneous in complexity. To illustrate this heterogeneity, in Figure 3 we present the distribution of the response time values for the PHP requests during the execution of the trace shown in Figure 2. The results were obtained using a fixed number of 3 (?) PHP servers, and show an irregular dispersion of the response time. The main reason for this dispersion is that the Wikipedia articles vary a lot in complexity, requiring different amounts of information and of items that need to be retrieved from the database.

Figure 2: Wikipedia trace workload.

Figure 3: Complexity of PhP requests.

Similarly, as depicted by Figure 2 and Figure 3, there is not any correlation between the PhP request volumes and their response times. More precisely in Figure 3, the highest response time values obtained in the interval of time 600min. match up with a drop in the request rate during the same interval in Figure 2. Therefore, any provisioning technique that makes decisions based on the request rate can incur errors by under- or over-provisioning an application, which can reduces the efficiency of the scaling actions.

In addition, other properties may be also considered important when using the Wikipedia traces such as the interarrival time between requests, the distribution of page popularity, the mix of static/dynamic requests, the ratio of read/write operations and requests for non-existing pages and files.

The workload-mix and the variable amount of data and visitors make of Wikipedia a valid example of elastic web application. In this paper, we focus in the scalability of the PhP web hosting service, and thereby as the number of PhP servers hosting MediaWiki scale out or back based on the demanding workload.

Resource provisioning algorithms

In this section, we describe the different provisioning techniques implemented in ConPaaS.

Trigger-based provisioning

The existing cloud infrastructures provide means to implement provisioning mechanisms that adjust the amount of allocated resources based on a number of standard monitoring parameters. These are simple trigger-based systems that define threshold rules to increase or decrease the amount of computational resources when certain conditions are met. As an example, the Auto Scaling system offered by Amazon EC2 [2] can be used to define rules for scaling out or back an application based on monitoring parameters like CPU utilization, network traffic volume etc. This trigger-based technique is currently used in well-known cloud platforms such as RightScale [15] or OpenShift [13].

For the sake of comparison, we decided to design and implement a trigger-based provisioning mechanism in ConPaaS. This technique monitors CPU usage and application response time metrics, and dynamically adjusts the computational power allocated

to an application by analyzing whether the monitoring data exceed their thresholds [1], as illustrated in Algorithm 1. Obviously, the lower and upper bound of each threshold are pre-defined by the user before execution.

Algorithm 1: Load-based

Data:
    System-level metrics (CPU, Resp. time)
      - Pre-defined metric threshold ranges, thr
Result: Scaling decisions

while *auto-scaling is ON* **do**
  | Collect monitoring data of each metric, $data$;
  |
  | **if** *no recent scaling operation* **then**
  | | **if** $avg(data_i) >=$ threshold_min$_i$ **then**
  | | | ADD resources;
  | | **else if** $avg(data_i) <$ threshold_max$_i$ **then**
  | | | REMOVE resources;
  | | **end**
  | **end**
  | Sleep for 5 minutes ;
end

Even though these type of mechanisms are simple and widely used in cloud platforms, they are excessively reactive and not so precise when provisioning web applications due to several factors:

- **Workload heterogeneity:** In some web applications, the system performance behavior fluctuates following an irregular pattern caused by the heterogeneity of the workload mix, which makes it difficult to predict the system's behaviour.

- **Reactiveness:** An excessively reactive algorithm can affect the system's stability, by causing frequent fluctuations in the number of allocated resources; this has negative effects on the performance, as well as increases the infrastructure cost.

- **Resources heterogeneity:** The performance of virtual instances provided by current clouds is largely heterogeneous, even among instances of the same type, as shown in [3]. Simple trigger-based provisioning systems do not take this heterogeneity into account, thus providing less efficient resource allocation.

---

[1]Every metric has a threshold range (max,min).

Based on these factors, we believe trigger-based provisioning mechanisms can be improved without drastically increasing their complexity. A possible solution is the utilization of techniques that handle workload and resource heterogeneity without being excessively reactive. Moreover, the implementation of these techniques should remain simple to facilitate their integration in existing auto-scaling systems. In the following we present two techniques that aims at solving the aforementioned drawbacks by relying on predictive and more accurate methods.

Feedback provisioning

CORINA: *Generic suggestion about this section: I think we should give some better motivation/intuition about the 3 techniques presented in this section. Ideally it would have been to have separate experimental results with each technique used just on its own, and see how much improvement each technique brings. But since we don't have that, maybe you could just give some more explanations/intuition based on what you saw in the experiments, of why each technique works?*

HECTOR: *You mean to completely change the Section 3 and 4. I don't know if this will delay a lot the publication of this article.*

Based on our previous knowledge from load-based provisioning, we designed and implemented an algorithm which improves the accuracy of our scaling actions when hosting web applications. To achieve that, our algorithm relies on three simple mechanisms: the definition of weights to each metric included in the performance requirements, the use of flexible thresholds and the estimation of the workload trend.

CORINA: *I think the focus of this paragraph should not be to explain why we use multiple metrics instead of one (because the naive algorithm also uses multiple metrics) but why there are weights assigned to the metrics. Does this bring better performance/accuracy? Also another possible advantage of weights is that it can make the algorithm more generic – e.g. we can use another set of metrics for*

other applications and let the user configure the weights.

Weighted metrics: Through the definition of weight values to application and system metrics, our algorithm takes into consideration its weight when making scaling decisions. This mechanism allows to improve the accuracy of our decisions, as some metrics can rapidly alert of the existence of any degradation before than others. As an example, when hosting web applications, our algorithm associates weights in an ascending order to the following metrics: request rate, CPU usage and response time. Accordingly the response time has a higher weight than the request rate, since higher values in the response time rapidly indicate the existence of a performance degradation in a web application.

Algorithm 2: Feedback

Data:
    System and App-level metrics
      - Pre-defined metric threshold ranges, *threshold*
    Define the weight of each metric, *w*
Result: Scaling decisions

Create a queue to store historical workload, q;
Establish flexible thresholds from *threshold* ranges;
    - Scalings operations could be triggered, *pred_thr*;
    - Scalings operations must be triggered, *reac_thr*;

while *auto-scaling is ON* **do**
    Collect monitoring data of each metric, *data*;
    if *avg(data$_i$) >= pred_thr$_i$* **then**
      | Increment chances of *scaling_out* using w_i, *s_out*;
    else if *avg(data$_i$) < pred_thr$_i$* **then**
      | Increment chances of *scaling_in* using w_i, *s_in*;
      else
        | Decrease the values of *s_in* or *s_out*;
      end
    end

    Add to q the most recent workload value;
    Estimate historical workload trends (last ∼30min), *td*;
      - If trend is increasing then *td = 1*;
      - If trend is decreasing then *td = 0*;
      - Undetermined *td = -1*;

    if *no recent scaling operation* **then**
      if *avg(data$_i$) >= reac_thr$_i$ and td = 1 and s_out > s_in* **then**
      | ADD resources;
      else if *avg(data$_i$) < reac_thr$_i$ and td = 0 and s_out < s_in* **then**
      | REMOVE resources;
      end
      Reset *s_in* and *s_out*;
    end
    else
    | Sleep during 5 minutes;
    end
end

CORINA: *I think here there should be explained better the usage and intuition behind s_in and s_out. They are used as extra conditions for scaling out and back, but I'm not sure I understand exactly when they play a role: if the response time exceeds the reactive threshold, in what situations can the extra condition with s_in and s_out prevent scaling out? Also, are s_in and s_out ever reset or decremented?*

Flexible thresholds: This algorithm uses two levels of threshold ranges for each metric: *predictive* and *reactive*. As shown on Figure 4, there are two "head rooms" between the SLO (Service Level Objective) threshold (performance requirements predefined by the user) and the flexible thresholds. First, the head-room H_1 is between the predictive threshold and the reactive thresholds, and is intended to alert of possible workload alterations in advance. Thereby, when the system performance exceeds the predictive ranges, there is an increment in the chances of scaling actions will be triggered to tackle future SLO violations. Otherwise, there is a decrease of the scaling chances. In Algorithm 4, the variables *s_in* and *s_out* represents the scaling changes. Note that, the values of *s_in* and *s_out* are calculated based on the metric's weight. The second head-room H_2 comprises between the SLO and reactive thresholds is used to trigger scaling actions. Performance values exceeding the reactive threshold launch scaling operations if other conditions are also satisfied. As an example of flexible thresholds for the CPU-usage, we established a predictive range comprised between 30% and 70% , and a reactive comprised between 20% and 80%.

CORINA: *Are the terms "constant" and "temporal" alteration normally used in trend estimation, or are they terms that you defined yourself? In the second case I think we should mention something like "we define constant alteration as...". And about flash crowds: from what I know, flash crowds last usually longer (in the order of hours) and we do want the system to provision extra resources in such a case. So if there is a spike of high workload that lasts for 10-15 minutes, I don't think we should call it a flash crowd – maybe just spike*
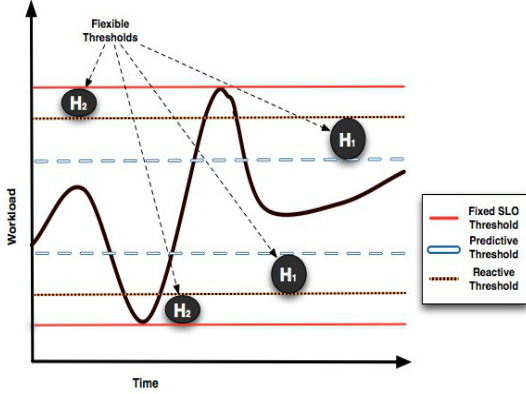
Figure 4: Flexible thresholds

or fluctuation or something like this. **Workload's trend estimation:** Nowadays, there is a wide literature on mathematical models that try to predict future alterations in web application's workload. However, the workload mix and network traffic of web applications make more difficult to provide accurate predictions using these models. Besides, the complexity of these models sometimes prevent its integration into real auto scaling systems. To design a robust and simple provisioning system, we decided to use a feedback mechanism that analyzes the behavior of the system performance during an interval of time. An exhaustive analysis of the monitoring data during a small interval of time (approx. during the last 30min) provides enough information to detect the workload's trend, and thereby to classify the type of workload alteration as *constant* or *temporal*. Obviously, only *constant* variations may trigger scaling actions to avoid frequent fluctuations in the system performance caused by short and sudden *temporal* variations (traffic spikes), as we will detail in Section . Note that, the workload trend estimation is not the trigger of this algorithm, but one more mechanism in it. So, it could be also calculated by using mathematical model such linear regression, kernel canonical correlation, and so on.

Anyway, the use of these three mechanisms must follow an order when making scaling decisions, as il-

lustrated in Algorithm 2. Initially, the user has to specify the thresholds ranges and give a weight to each metric. Next the flexible thresholds are defined based on Amazon recommendations and statistically-chose performance measures [2]. Once the monitoring data is collected from the agents, the decision making process can start. Firstly, these data is analyzed to verify if it exceeds the predictive threshold ranges (denoted by p_thr), if so the probability of triggering scaling actions increases proportionally in function of the metric's weight. In order to keep track of workload variations, this algorithm stores in a queue (denoted by q) the most recent system performance values, and analyzed them to estimate the workload trend (denoted by td). Finally, to trigger any scaling action, a serie of conditions have to be satisfied: (i) no previous scaling actions have been taken over the last 15min; (ii) the recent monitoring data have to exceed the predictive and reactive threshold ranges; (iii) the workload trend has to follow a constant pattern (increasing/decreasing).

Although the combination of these techniques improves the accuracy of our measurements, and avoids to present an excessive reactive behavior. The heterogeneous nature of the VM instances and the workload require more flexible provisioning algorithms, as we pointed out in [9].

Dynamic load balancing weights

Another problem that we need to consider is the heterogeneity of cloud platforms. Different virtual machines from the same cloud might have different performance characteristics, even when their specifications from the cloud vendor are the same. This problem can be addressed through various load balancing techniques, like assigning weights to the backend servers or taking into account the current number of connections that each server handles. Furthermore, the performance behavior of the virtual servers can also change in time, either due to changes in the application's usage patterns, or due to changes related

---

[2]These performance values are obtained from previous executions of the same application using similar hardware configurations.

to the hosting of the virtual servers (e.g., VM migration).

In order to address these issues in ConPaaS we implemented a weighted load balancing system in which the weights of the servers are periodically re-adjusted automatically, based on the monitoring data. As illustrated in Algorithm 3, this method assigns the same weight to each backend server at the beginning of the process. The weights are then periodically adjusted (every $\sim$ 15min) proportionally with the difference among the average response times of the servers during this time interval.

Algorithm 3: Dynamic load-balancing weights

Data:
    System and App-level metrics
      - Pre-defined metric threshold ranges, *threshold*
    Define the weight of each metric, *w*
Result: Scaling decisions

Create a queue to store historical workload, q;
Establish flexible thresholds from *threshold* ranges;
    - Scalings operations could be triggered, *pred_thr*;
    - Scalings operations must be triggered, *reac_thr*;
Initialize load-balancing weights for the backend servers;

while *auto-scaling is ON* **do**
  Collect monitoring data of each metric, *data*;
  if $avg(data_i) >= pred\_thr_i$ **then**
  | Increment chances of *scaling_out* using w_i, *s_out*;
  else if $avg(data_i) < pred\_thr_i$ **then**
    Increment chances of *scaling_in* using w_i, *s_in*;
    else
    | Decrease the values of *s_in* or *s_out*;
    end
  end

  Add to q the most recent workload value;
  Estimate historical workload trends (last $\sim$20min), *td*;
    - If trend is increasing then $td = 1$;
    - If trend is decreasing then $td = 0$;

  if *no recent scaling operation* **then**
    if $avg(data_i) >= reac\_thr_i$ *and* $td = 1$ *and* s_out $>=$ s_in **then**
    | ADD resources;
    else if $avg(data_i) < reac\_thr_i$ *and* td = 0 *and* s_out $<$ s_in **then**
    | REMOVE resources;
    end
    Reset *s_in* and *s_out*;
  end
  else
  | Sleep during 5 minutes;
  end

  Adjust load-balancing weights based on the workload ($\sim$15min);
end

Evaluation

In this section we conducted our experiments on a heterogeneous infrastructure like Amazon EC2 [2], and on a homogeneous infrastructure like DAS-4 (the Distributed ASCI Supercomputer 4) [1]. DAS-4 is the Dutch Computational Infrastructure, a six-cluster wide-area distributed system designed with research purposes. In our experiment campaign, we compared the degree of SLO enforcement (performance requirements fulfillment) and resource consumption for each one of the provisioning algorithms included in ConPaaS.

Testbed configuration

As a realistic and representative scenario, we deployed MediaWiki application using ConPaaS on both infrastructures, and we ran the Wikibench tools utilizing Wikipedia workload traces.

Our goal is to evaluate the behavior of the provisioning algorithms, when scaling out and back the number of VMs hosting PhP servers to guarantee several performance requirements, referred to as SLO. Specifically, we fixed two SLOs, one of 700 milliseconds at the service's side (denoted by a yellow Line on Figures 5,6,8,9 and 10), and another of 1500 milliseconds at the client's side (denoted by a red Line on Figures 5,6,8,9 and 10). Thus, our measurements shows the behavior of the MediaWiki application under a workload generated from real access traces during 24h. Note that, these experiments only focus on the average of PhP response times and the resource consumption obtained with our algorithms. Accordingly, some assumptions were made:

- Response times from static requests were not analyzed due to its lightweight nature.

- The algorithms collect the monitoring data through Ganglia over a reporting period of 5 minutes.

- Since DAS-4 is a homogeneous infrastructure, the dynamic load-balancing weights provisioning technique was only evaluated on a heterogeneous platform like Amazon EC2. The benefits behind

of this algorithm can only be appreciated when running on a heterogeneous environment, that provides VMs with different hardware configurations.

- These experiments used the same statistically-chose performance threshold for both infrastructures. In the future, these threshold values might be defined based on the type of VM provisioned.

- A minimum interval of 15 minutes has been established between scaling actions.

Homogeneous Infrastructure

Our experiments on DAS4 relies on OpenNebula as Infrastructure-as-a-service (IaaS) [16]. To deploy the Wikipedia services, we used small instances for the PhP service (manager and agents) and a medium instance for the MySQL service (agent). OpenNebula's small instances provision VMs equipped with 1 CPU of 2Ghz, and 1GiB of memory, while medium instances are equipped with 4 CPU's of 2Ghz, and 4GB of memory.

SLO enforcement. Figure 5 and Figure 6 depict the degree of SLO fulfillment of the trigger-based and feedback algorithms, indicating the average of response times obtained during the execution of the Wikipedia workload trace. As illustrated on Figure 5, the results show how the trigger-based provisioning algorithm presents an important amount of SLO violations (between the yellow and red Lines), which are generated due to its excessive reactive behavior. As we mentioned, this algorithm is an easy target to traffic spikes, as it tends to add or remove VMs to handle sharp and sudden variations in the workload, thus producing SLO violations. In contrast on Figure 6, the system performance (*i.e.,* response time) do not fluctuate greatly showing a more stable behavior during the whole experiment. As a result, there is a drop of the 31.72% in the amount of SLO violations in comparison with the trigger-based algorithm (when regarding the SLO's at the client side).


Figure 5: PhP resp. time on DAS4 – Trigger-based.

Resource consumption. Nevertheless to better understand the behavior of both algorithms, we may also focus on the resource consumption illustrated on Figure 7. Firstly, the excessive reactive behavior of the trigger-based algorithm is again illustrated at the interval $t=350min$ and $t=820min$, where two scaling operations under-provision the system during a short period of time. These provisioning decisions provoked fluctuations in the system performance that incremented the financial cost, as well as throughput alterations under the same intervals of time, as depicted on Figure 5. When using the feedback algorithm, the system makes provisioning decisions by analyzing workload's trend during a considerable interval of time. Scaling actions are only triggered when having constant alterations in the Wikipedia workload, thereby providing a more efficient resource usage. Indeed, the workload alterations depicted on Figure 2, match with the provisioning decisions made by the feedback algorithm on Figure 7.

Discussion. Using the trigger-based provisioning algorithm, the system performance fluctuates greatly following a pattern similar to the web traffic, that increases the number of SLO violations. The reactive behavior of this algorithm triggers scaling actions that affect to the system performance instead of im-

9

Figure 6: PhP resp. time on DAS4 – Feedback.

Figure 7: Resource consumption on DAS4.

proving it, and as a consequence, it is also wasteful in terms of resource consumption. Unlike feedback algorithm offers an efficient resource usage and a constant performance behavior while meeting the application's SLO.

Both algorithms are best-effort regarding the SLO fulfillment, and thereby temporal alterations of the workload (with a short duration of 5min approx.) cannot be handled. The heterogeneity of the PhP-requests including images and requiring multiple Db queries, as well as the startup time of VMs (2-5min), are in part responsible of these SLO violations.

Heterogeneous Infrastructure

Our experiments on EC2 used small instances for the PhP service (manager and agents) and a medium instance for the MySQL service (agent). EC2 small instances provision VMs equipped with 1 EC2 CPU, and 1.7GiB of memory, while medium instances are equipped with 2 EC2 CPU's, and 3.75GiB of memory.

SLO enforcement. Figure 8, Figure 9 and Figure 10 show the system performance of the trigger-based, feedback and dynamic load-balancing weights algorithms, respectively. As depicted on Figure 8, the performance fluctuates greatly following an ir-

regular pattern when using the trigger-based algorithm. More precisely, two of the three workload peaks caused at $t=300min$ and $t=820min$, are explained by looking at the variations on the Wikipedia workload described on Figure 2. However, there is a third peak between $t=400min$ and $t=500min$ that corresponds to the interval of time on which the workload trace experiences a significant drop in the request volumes. During this period of time, the workload-mix causes sudden and sharp performance variations that explains this behavior, as a side effect associated to very frequent scaling actions. As a result, there is a degradation of the SLO fulfillment.

On the other hand, Figure 9 and Figure 10 show as the feedback and dynamic load-balancing weight algorithm behave similarly. Even though both algorithms are best-effort, there is an important reduction in the number of SLO violations during the trace execution. In particular, the feedback algorithm reduces the SLO violations in a 41.3%, while the dynamic load-balancing weights algorithm does it in a 47.6%. Like on DAS-4, the feedback algorithm follows a constant performance pattern without having sharp and sudden workload alterations. Besides, as shown on Figure 10, the dynamic load-balancing weights algorithm algorithm has a similar behavior to the feedback algorithm in terms of system performance, however. This algorithm improves the SLO enforcement in a 6.3% in comparison with the feedback algorithm

Figure 8: PhP response time on EC2 – Trigger-based.

Figure 9: PhP resp. time on EC2– Feedback.

at the client's side. Therefore we demonstrate how the use of workload-mix and flexible load-balancing techniques, although intrusive, do not cause time delays or excessive throughput alterations.

Resource consumption. The resource usage on EC2 presents important alterations, as shown on Figure 11. When using the trigger-based provisioning, the fluctuations in the system performance are explained as a result of a high frequency of scaling operations. In concrete, the fluctuations caused at the interval of time between $t=400min$ and $t=500min$ (see on Figure 8), match with the provisioning decisions made during the same interval of time on Figure 11. If we now pay attention to the feedback, and dynamic load-balancing weights algorithms, their resource consumptions are identical along the execution. Indeed, both algorithms decided to scale out the system during the interval of time comprised between $t=1050min$ and $t=1400min$, to prevent future SLO violations that occurred when using the trigger-based algorithm. In particular, this situation demonstrates the benefits of using flexible threshold ranges to provide a predictive provisioning mechanism, thus improving the user experience.

Discussion. The experiments on EC2 leads to several conclusions. The use of the trigger-based algorithm shows again how aggressive provisioning increases the resource consumption and the chances of degraded application performance due to frequents scaling actions. These actions are triggered as an effect associated to the workload-mix, when handling bursty workload conditions. On the other hand, the feedback and dynamic load-balancing algorithms constitute two robust provisioning models that offer an efficient resource consumption and keep stable the application performance during the trace execution.

Furthermore, the use of a dynamic load-balancing algorithm provided a more efficient distribution of the request-mix across servers, that reduced the SLO violations in a 6.3%. Hence, the main objective behind of this algorithm is to tackle the degradations caused by the workload mix.

Discussion

Generally, the result of our measurements show how the behavioral performance pattern and the resource consumption vary depending on the infrastructures on which we ran our experiments. Different hardware configurations such as those provided by DAS-4 and EC2, offer two distinct scenarios to validate our

11

Figure 10: PhP resp. time on EC2– Load-balancing Weights.

Figure 11: Resource consumption on EC2.

provisioning algorithms. In these experiments, we demonstrate how trigger-based provisioning mechanisms can affect the system performance instead of improving it, as well as are wasteful in terms of resource usage. Furthermore, we show how a dynamic load-balancing technique, although intrusive, can be included and used without producing performance alterations. In fact, this technique slightly reduced the number of SLO violations in comparison with the results obtained using feedback algorithm. Finally, we also present the benefits by using feedback and dynamic load-balancing weights provisioning algorithms which aims to find the trade-off between the accuracy and cost savings.

However, there is room for improvement using the dynamic load-balancing weights algorithm, as some workload alterations could not be handled during the trace execution.

Related studies

There is a wide literature on issues related to dynamic resource provisioning for cloud web applications. Different approaches present solutions based on queuing models [18], feedback loops techniques [7], mathematical models [12] or even approaches using neural networks techniques [8]. However, most of these models require a deep understanding in mathematics or machine learning techniques which are not easily interpreted by non specialists. Besides the traffic in web applications is shaped by a combination of different factors such as diurnal and seasonal cycles, sociological and psychological, that follows an irregular pattern. It makes extremely challenging the design and development of realistic and accurate dynamic provisioning mechanisms.

These well-known drawbacks force to IaaS like Amazon EC2 and Windows Azure [11], or PaaS like RightScale [15] and OpenShift [13], to design simple threshold rule-based auto-scaling systems, instead of relying on approaches from academic research. Unfortunately, these scaling systems are naive, wasteful in terms of resource consumption and cost savings, and an easy target for flash crowds.

In the following, we present some of the most relevant and realistic academic approaches that proposed dynamic resource provisioning mechanisms for multi-tier applications.

In [17, 18], the authors designed and implemented a predictive and reactive provisioning mechanism. They used a queuing model G/G/1 to decide the server pool size to be provisioned, and an admission control mechanism to face extreme workload variations. Offline prolifing techniques were employed to

gather information about the resource requirements of the incoming requests for each tier, and thereby to selectively admit/reject requests for the lightweight files. An evaluation using real-traces on a homogeneous infrastructures shows the benefits of this approach when handling flash crowds. Unfortunately, its admission control mechanism incurs into sporadic SLA violations (if the server utilization exceed a predefined threshold) reducing the QoS of the service, and therefore affecting user experience. Similarly to the previous work, [19] extended queuing models and transaction mix models to design a predictive and reactive provisioning system. To model the application performance, they integrated proactive control and feedback control methods that dynamically adjusted the CPU capacity allocated to servers. This work only considered SLA constraints at the system-level, while others constraints at application-specific level such as response time and request rate were not taken into consideration. Besides, an evaluation of CPU variations on a homogeneous infrastructure, when processing traces from a non real-world application, lack arguments to valid its approach.

Regarding the management of flash crowds [21], a proactive application workload manager was designed to separate the user requests between two groups of servers: one named 'base workload' referred to the smaller and smoother variations in the workload; and the other 'tresspassing' referred to the temporal burstly workloads caused by flash crowds. To do this, the authors attempt to divide the data items into popular and less popular, and place them in the right group of servers. Even tough a realistic evaluation was conducted on Amazon EC2 utilizing real traces (Yahoo video streaming), authors do not explain in details how the dynamic resource provisioning is done. Recently, in [10], online profiling techniques have been utilized for managing the tradeofff between performance overload, and cost savings for dynamic resource provisioning. The authors replicate at runtime a regular server hosting an application, with a new server with profiling instrumentation. Their experimental results show how profiling techniques can be included in a resource provisioning system, without causing important response time delays or throughput alterations in comparison with non-profiling provisioning. As we mentioned in Section , profiling techniques can report more benefits than performance degradations or expenses.

Acknowledgments

> **FIXME: Probably also mention ERRIC here. ??**

References

[1] Advanced School for Computing and Imaging (ASCI). The Distributed ASCI SuperComputer 4. `http://www.cs.vu.nl/das4/`.

[2] Amazon Elastic Compute Cloud (EC2). `http://aws.amazon.com/ec2`.

[3] Jiang Dejun, Guillaume Pierre, and Chi-Hung Chi. EC2 performance analysis for resource provisioning of service-oriented applications. In *Proceedings of the 3rd Workshop on Non-Functional Properties and SLA Management in Service-Oriented Computing*, November 2009.

[4] Anh Vu Do, Junliang Chen, Chen Wang, Young Choon Lee, A.Y. Zomaya, and Bing Bing Zhou. Profiling applications for virtual machine placement in clouds. In *2011 IEEE International Conference on Cloud Computing (CLOUD)*, pages 660–667, July 2011.

[5] Guido Urdaneta Erik-Jan van Baaren and Guillaume Pierre. Wikibench: A distributed, wikipedia based web application benchmark. 2009.

[6] Ganglia monitoring system. `http://ganglia.sourceforge.net/`.

[7] Jiayu Gong and Cheng-Zhong Xu. A gray-box feedback control approach for system-level

peak power management. In *2010 39th International Conference on Parallel Processing (ICPP)*, pages 555 –564, September 2010.

[8] Sadeka Islam, Jacky Keung, Kevin Lee, and Anna Liu. Empirical prediction models for adaptive resource provisioning in the cloud. *Future Gener. Comput. Syst.*, 28(1):155–162, January 2012.

[9] Dejun Jiang. *Performance Guarantees For Web Applications*. Phd thesis, VU University Amsterdam, March 2012.

[10] Nima Kaviani, Eric Wohlstadter, and Rodger Lea. Profiling-as-a-service: adaptive scalable resource profiling for the cloud in the cloud. In *Proceedings of the 9th international conference on Service-Oriented Computing*, ICSOC'11, pages 157–171, Berlin, Heidelberg, 2011. Springer-Verlag.

[11] Microsoft Windows- Windows Azure. `http://www.windowsazure.com/en-us`.

[12] Sireesha Muppala, Xiaobo Zhou, Liqiang Zhang, and Guihai Chen. Regression-based resource provisioning for session slowdown guarantee in multi-tier internet servers. *Journal of Parallel and Distributed Computing*, 72(3):362–375, March 2012.

[13] OpenShift. `https://openshift.redhat.com/app/flex`.

[14] Guillaume Pierre and Corina Stratan. ConPaaS: a platform for hosting elastic cloud applications. *IEEE Internet Computing*, 16(5):88–92, September-October 2012.

[15] RightScale. `http://www.rightscale.com//`.

[16] B. Sotomayor, R.S. Montero, I.M. Llorente, and I. Foster. Virtual infrastructure management in private and hybrid clouds. *IEEE Internet Computing*, 13(5):14 –22, October 2009.

[17] Bhuvan Urgaonkar and Prashant Shenoy. Cataclysm: Scalable overload policing for internet applications. *Journal of Network and Computer Applications*, 31(4):891–920, November 2008.

[18] Bhuvan Urgaonkar, Prashant Shenoy, Abhishek Chandra, Pawan Goyal, and Timothy Wood. Agile dynamic provisioning of multi-tier internet applications. *ACM Trans. Auton. Adapt. Syst.*, 3(1):1–39, March 2008.

[19] Zhikui Wang, Yuan Chen, D. Gmach, S. Singhal, B.J. Watson, W. Rivera, Xiaoyun Zhu, and C.D. Hyser. AppRAISE: application-level performance management in virtualized server environments. *IEEE Transactions on Network and Service Management*, 6(4):240 –254, December 2009.

[20] Wikipedia Foundation. MediaWiki a free software wiki package. `http://www.mediawiki.org`.

[21] Hui Zhang, Guofei Jiang, Kenji Yoshihira, Haifeng Chen, and Akhilesh Saxena. Resilient workload manager: taming bursty workload of scaling internet applications. In *Proceedings of the 6th international conference on Autonomic computing*, ICAC '09, pages 45–46, New York, NY, USA, 2009. ACM.