

Practical Performance Control for Web Applications in the Cloud

Corina Stratan, Guillaume Pierre, Hector Fernandez, Eliana Tirsia and Valentin Cristea

Vrije Universiteit Amsterdam and University Politehnica of Bucharest

Other versions for the title:

Automatic Scaling for Cloud Web Applications in Practice

A Practical Approach for Automatic Scaling (Resource Provisioning?) of Web Applications in the Cloud

Implementing Automatic Scaling for Cloud Web Applications ;

1 Introduction

Over the last few years, we have witnessed numerous debates on blogs, conferences and media on whether to use *grid computing* or *cloud computing*. Even though some of the early promises of grid computing, sound the same as the promises of the cloud, cloud efforts were focused on applications which will be long-running on the infrastructure. As a result, today's most well-known enterprises use cloud infrastructures as a tool on which deploy their applications. Similarly, one key properties of the clouds is *elasticity*, *i.e.*, the possibility for cloud users, to dynamically adapt the quantity of resources at his disposal at run-time. Unfortunately, this notion of infinite resource availability has a financial cost, as cloud users have to pay for the resources being used from moment to moment.

CORINA: *Not sure if we should bring grid computing into discussion, the readers might not be familiar with it (and I think we should aim to shorten the paper).*

Therefore it seems crucial to choose the best cloud provisioning offers in order to minimize the resource consumption, and thereby the running costs while assuring several performance requirements for our

applications. Nowadays, there are lots of research papers dedicated to sophisticated techniques on dynamic resource provisioning. However, if we look at real deployments we see that cloud providers rely on extremely simple techniques, and completely ignore the results from academic research on the topic. Two reasons can be the explanations for this discrepancy: *(i)* the gains of using sophisticated provisioning strategies are too low for anyone to bother; *(ii)* implementing these techniques is a difficult exercise, which is why real cloud systems rely on simpler techniques.

CORINA: *I think the statement “completely ignore the results from academic research” is a bit too bold. I would say something like the real deployments tend to use very simple strategies rather than the sophisticated techniques proposed by researchers.*

This paper tries to identify the real cause. We do this by implementing a sophisticated provisioning system in realistic conditions, and reporting on *(i)* how hard implementation was; and *(ii)* potential gains from using the better technique as compared to a simple strawman. To achieve this, we designed and implemented these resource provisioning strategies on ConPaaS, an open source platform-as-a-service environment for hosting cloud applications [13]. For the sake of comparison and discussion, we deployed the Wikipedia services and used real access traces to validate our technique, opening doors to real implementation of slightly sophisticated auto-scaling systems.

FIXME: Are we gonna include the three strategies?

CORINA: *Maybe not; we could just make a summary phrase out of that (e.g. we*

used more sophisticated methods based on trend detection, dynamic weight adjustment for load balancing etc.).

FIXME: Do you have any preference about terminology: strategy, method, algorithm or technique

CORINA: *I think these terms can be used interchangeably, I don't have a preference.*

Section 2 introduces the ConPaaS runtime environment. Section 3 describes the application and its realistic benchmark to validate our system. Section 4 focus on the different resource provisioning strategies implemented in ConPaaS. Section 5 details the experimental campaign and its results. Section 6 discusses related works. Section 7 draws a conclusion.

2 ConPaaS overview

ConPaaS is an open-source runtime environment for hosting applications in Cloud infrastructures. Within the Cloud computing paradigm, ConPaaS belongs to the platform-as-a-service family, in which a variety of systems aim to simplify the deployment of applications in the Cloud. Using ConPaaS, developers can now focus their attention on application-specific concerns rather than making their applications more suitable for the cloud.

2.1 Architecture

In ConPaaS, an application is conceived as a composition of one or more elastic and distributed components, called *services*. Each service is dedicated to host a particular type of functionality of an application. At the moment, ConPaaS supports six different types of services: two web application hosting services respectively specialized for hosting PHP and JSP applications; a MySQL database service; a NoSQL database service built around the Scalarix key-value store; a MapReduce service; and a Task-Farming service for high-performance batch processing.

These services are built based on an architecture composed of two main building blocks: agents and managers.

CORINA: *Maybe include here one of the diagrams that we have, that show the architecture of ConPaaS?*

- **Agent:** A service is composed of one or several agents VMs which host the needed components to provide the service-specific functionality. As an example, the PhP web hosting service needs at least three agents: one http proxy, one web server and one PhP server.
- **Manager:** For each service, there is only one manager VM. The manager is in charge of executing all management requests, centralizing governance and performance monitoring data, and controlling the allocation of resources assigned to one service. The actual application traffic is not addressed to the manager. Requests from end users willing to access the application must be directed directly to the agents hosting the application.

CORINA: *We should point out that the agents can be on a single VM (for small applications) or on multiple VMs (for larger applications).*

2.2 Hosting Elastic Applications

The main features that distinguish ConPaaS from other PaaS systems are its approach for autonomous application scaling and its interoperability with a wide variety of private and public IaaS clouds. In particular to provide such autonomous scaling capabilities, ConPaaS includes a monitoring data analysis mechanism and a resource provisioning system.

Since most IaaS does not implement any sophisticated monitoring capabilities, ConPaaS incorporates a scalable distributed monitoring engine which is based on the Ganglia [5] monitoring system. Ganglia consists of a server component (gmetad) that aggregates monitoring statistics from various VMs, and a reporting agent (gmond) which runs inside each VM. By default, Ganglia monitors only system-level metrics such as disk, CPU, memory and network usage. Unfortunately, these metrics often do not provide enough information about system performance due to the heterogeneity of the applications. As a

consequence, in ConPaaS, we enhanced ganglia to also monitor service workloads by enhancing the reporting agent to track service-specific logs at runtime, and report statistics over a reporting period of 5 minutes. For instance, the PHP web hosting service includes new ganglia metrics that report statistics about the response time and request rate for static and non/static user requests, respectively. Once the monitoring data is collected from the agents VMs, the resource provisioning algorithm decides whether to trigger scaling operations based on this data.

Unlike of implementing traditional trigger-based provisioning systems that scale services independently of whether they are part of an application. In ConPaaS, we designed a performance control model for multi-service applications. Our model takes into consideration the fact that services may have different roles in an application, so that services collaborate to guarantee several performance requirements. It improves the effectiveness and accuracy of the scaling decisions. Indeed, this allows to rapidly detect performance bottlenecks in applications, and thereby to minimize the resource consumption.

FIXME: Do you want to mention profiling in this section?

CORINA: *I don't think so, since we didn't even really use profiling in the experiments that we show here.*

3 The Wikipedia workload

To evaluate the behavior and accuracy of ConPaaS when hosting web applications, we prepared a realistic and complex enough scenario to assess any PaaS. In particular, we deployed the Wikipedia web application called MediaWiki [18], and used a web hosting benchmark called WikiBench [4].

The architecture of the Wikipedia website uses a http-proxy, a http-web server, a database and one or more PHP servers. To deploy the Wikipedia services on ConPaaS, we composed two different services: the PHP web hosting and MySQL service. In the MySQL service, we installed a complete copy of English Wikipedia database which contains about 30GB in Wikipedia articles. In the PHP service,




Figure 1: Wikipedia trace workload

an initial configuration was composed of one Nginx HTTP-proxy, one Apache HTTP server, and one or more PHP (FastCGI Process Manager) servers. Each PHP server hosts the MediaWiki application which is the main component of this system.

In order to benchmark ConPaaS when hosting the Wikipedia services, we used the WikiBench research tool which generates realistic benchmarks with adaptable traffic properties. WikiBench has a number of advantages compared to the existing benchmark tools for web applications. First of all, WikiBench traces add a high degree of realism, since it is entirely based on the Wikipedia software and data. Indeed, the benchmark workloads are generated based on real access traces from the Wikimedia Foundation. These traces contain detailed traffic logs of requests made to Wikipedia by its users. Since the original Wikipedia traces can reach peaks of 50000 or 60000 reqs./secs, WikiBench uses the original 10% sample of these traces which can generate a workload up to about 5000 reqs./secs. As an example, in Figure 1, we show the workload of one trace as the number of request per minute during approximately 24h.

Even though we use a 10% of the real traces, they are very heterogeneous regarding the type of user requests, and thus explaining the irregular performance

pattern followed by web applications. We highlight the following properties that we consider important about these requests:

- The intervarrival time between requests follows a Poisson distribution.
- The distribution of page popularity varies from very popular pages to those following a Zipf distribution, and finally to a large part of them being accessed very infrequently.
- The ratio of static/non-static file requests presents a strong variation. Even though, most requests are for images, css stylesheets and other static files. The types of requests are related, as the result of a Wikipedia web page includes images, css stylesheets and other static files. In particular, most web pages contain PHP code to be compiled that increases significantly the response time. We identify the reasons of this increment and we want to retain: (i) PHP web pages often require DB queries; (ii) PHP pages vary in complexity so it is difficult to make predictions. By complexity, we mean the size and number of required database queries to compile a PHP page.
- The ratio of read/write operations vary having more reads than editions or creations of wiki pages.
- A considerable amount of requests for non-existing pages and files, which obviously add realism to the traffic.

Since Wikipedia has a variable amount of data and visitors, it represents a valid example of elastic web applications. In this paper, we focus in the scalability of the PHP web hosting service, and thereby as the number of PHP servers hosting MediaWiki scale out or back based on the demanding workload.

4 Resource provisioning methods

In this section, we describe the different provisioning methods implemented in ConPaaS.

4.1 Naive provisioning

CORINA: *Maybe we should find another name instead of “naive”, for example “simple” or “basic”?*

The existing infrastructures offer provisioning mechanism that adjust the amount of resources depending on the load of the currently allocated resources. Within this group of provisioning mechanisms, we include the simple trigger-based systems that define threshold rules to increase and decrease the computational power of an application in order to guarantee several performance requirements. As an example, the Auto Scaling system offered by Amazon EC2 [2] scales out an application whenever its CPU usage in the last 10min exceeds a specific threshold (Amazon recommends values from 70%). This rules-based method is currently used in mature cloud platforms such as RightScale [14] or OpenShift [12]. In ConPaaS, we also decided to offer a trigger-based provisioning mechanism, called “naive provisioning”, that monitors CPU usage and response time thresholds.

Even though these type of mechanisms are simple and widely used in cloud platforms, they are naive and excessive reactive due to several factors:

- When hosting a web application, the system performance behavior fluctuates following an irregular pattern caused by the web traffic (workload).
- Excessive reactive system can affect the system performance when handling flash crowds. A high frequency of scaling operations can provokes sharp and sudden fluctuations that affect the performance instead of improving it. So, it is particularly difficult to decide when to scale out or back.
- Services are handled as black boxes, the definition of threshold rules only covers system-level constraints such as response time and CPU usage. Therefore, when provisioning web applications, constraints such as request rate of static/non-static files may be also taken into consideration.

- The performance of virtual instances provided by current Clouds is largely heterogeneous, even when requesting the exact same type of instance each time [3].

Based on these factors, we believe trigger-based provisioning mechanisms can be improved without drastically increasing its complexity. A solution seems to be the utilization of techniques that handle web application workload, while its implementation remains a simple exercise. In the following we present two algorithms that aims at solving these drawbacks by relying on more predictive and dynamic methods.

CORINA: *What do you mean by “its implementation remains a simple exercise”?*

4.2 History-aware provisioning

Based on our previous knowledge from naive provisioning, we designed and implemented a more predictive and dynamic algorithm to handle temporal bursty workload. To achieve that, our algorithm relies on three simple techniques: weight-values for SLA constraints, flexible thresholds and estimation of the workload trend.

Weight-values: once the system collected the monitoring data from application-specific constraints such as the CPU usage, response time and request rate for both non-static/static files. Traditional algorithms would scale out and back whenever CPU usage or response time values exceed a beforehand defined threshold range. Nevertheless, through the definition of weight values, our algorithm gives different priorities to each constraint when making decisions. Considering web applications, our method associate weight-values in an ascending order to the following constraints: request rate, cpu usage and response time. Accordingly the response time has a higher weight value than the request rate, since a high value in the response time rapidly indicate the existence of a performance degradation in a web application. High values in the request rate cannot always indicate that an application is becoming overloaded, due to the large diversity in the complexity of the requests.

Flexible thresholds: The history-aware algorithm establishes two types of threshold ranges: *predictive* and *reactive*. A *predictive threshold* comprises a range of values that alert of possible workload alterations in advance. While a *reactive threshold* comprises a range of values which are used to trigger scaling actions. As an example of CPU flexible threshold, we established a predictive range comprised between 30% and 70% , and a reactive comprised between 20% and 80%.

Workload’s trend estimation: Nowadays, there is a wide literature on mathematical models that try to predict future alterations in web application’s workload. However, most models fail due to the irregular pattern of the web traffic, or its complexity to be utilized in real auto scaling systems. As a result, we decided to analyze the system performance behavior during an interval of time, thus avoiding flash crowds and slashdot effects. An exhaustive analysis of the monitoring data during a small interval of time (approx. during the last 20min) provides enough information to deduct the workload’s trend, and thereby to classify the type of workload alteration as constant or temporal. Obviously, only constant variations may trigger scaling actions to avoid fluctuations in the system performance, as we will detail in Section 5.

CORINA: *Are you sure that the mathematical models from the existing literature fail? The fact that the traffic is irregular makes it difficult to provide accurate predictions, but trend detection does not aim to make accurate predictions, so I think those models should still work.*

Using the history-aware algorithm we now collect application-specific metrics for our measurements, evaluate the evolution of the application workload avoiding an excessive reactive behavior, and handle threshold values that helps to predict possible SLA violations. Unfortunately, the heterogeneous nature of the VMs requires more dynamic provisioning algorithms.

Profiling-based provisioning

CORINA: *This section describes the profiling methods, which were not really used in the experiments that we show in the paper. We should probably describe instead the method that we actually used, which is dynamically adjusting the load balancing weights based on the monitoring data (without profiling).*

The heterogeneity of cloud platforms, and therefore, of their VMs affect to the accuracy of the provisioning decisions. VMs with better hardware configuration can sustain higher workload intensities. Based on that, each time ConPaaS requests a new VM from the Cloud, it must measure the individual performance profile of this particular VM instance before being able to decide what is the most efficient use of this resource. To do that, ConPaaS uses profiling techniques which represent a promisingly solution to this heterogeneity problem, as we pointed out in [8]. There are two types of performance profiling: offline and online.

- *Offline profiling* techniques gather information about VM instances to have an initial assessment of their throughput. This assessment will be used to predict SLA violations in advance. For instance, offline profiling can be used to define flexible threshold ranges for each VM (*i.e.*, max. and min. request rate, etc...).
- *Online profiling* techniques enable to decide at runtime which set of VM instances is more suitable to handle the current workload intensity. Once the application started, online profiling can be also used to dynamically re-adjust the workload balance across VMs based on its own performance behavior. Specifically, the profiling-based algorithm uses a mechanism to continuously adapt the online profile, in order to deal with possible changes in workload's type or in VMs performance. This online profile is used to dynamically adapt the load balancing weights and also our performance predictions.

As illustrated on Figure 4.2, we believe most existing problems in resource provisioning can be solved




Figure 2: Profiling Resource Provisioning

by combining the history-aware provisioning algorithm (referred as scaling on the Figure 4.2) in conjunction with these two profiling techniques. The resulting provisioning system scales out/back an application utilizing all the benefits of heterogeneous cloud infrastructures in a predictive and dynamic manner.

5 Experimental evaluation

In this section we conducted our experiments on a public cloud like Amazon EC2 [2] and on a private cloud like DAS-4 (the Distributed ASCI Supercomputer 4) [1]. DAS-4 is the Dutch Computational Infrastructure, a six-cluster wide-area distributed system designed with research purposes. We compare the degree of SLA fulfillment and resource consumption for each one of the provisioning algorithms included in ConPaaS.

Note that, since DAS-4 is a homogeneous infrastructure, the profiling techniques were only evaluated on a heterogeneous platform like Amazon EC2.

FIXME: Public cloud or Homogeneous infrastructure ??

CORINA: *I would call it homogeneous infrastructure. The fact that the cloud is public or not doesn't have any influence on our experiments (and I would say EC2 is also a public cloud).*

FIXME: Include all the assumptions ??

CORINA: *I guess only the important ones (but what kind of as-*

assumptions were you thinking about?

FIXME: I did not mention the size of our checking window: 5min, 5min also

CORINA: *I think we should mention that.*

5.1 Testbed configuration

As a realistic and representative scenario, we deployed Mediawiki application using ConPaaS on both clouds, and we ran the Wikibench tools utilizing Wikipedia workload traces.

To provide the Wikipedia services, an initial configuration was composed of 4 VMs, and 1 VM to host the Wikibench tools. The 4 VMs include a PHP service manager, a FPM-PHP agent, a web server and a http-proxy agent (in the same VM), and a MySQL service to store the English Wikipedia data, as explained in Section 3.

Thus, the provisioning system will scale out and back the number of VMs hosting FPM-PHP agents to guarantee the SLO (Service Level Objective). Initially, we fixed two SLOs one of 700ms (milliseconds) at the service’s side and another of 1500ms at the client’s side. Our measurements shows the behavior of the Wikipedia services during 24h under a workload generated from real access traces. Note that, our experiments only focus on the average of PHP response time and the resource consumption of our algorithms. The response time of static file requests is not evaluated due to the lightweight nature of the static files employed by Wikipedia articles.

5.2 Public Cloud

Our experiments on DAS4 relies on OpenNebula as Infrastructure-as-a-service (IaaS). To deploy the Wikipedia services, we used small instances for the PHP service (manager and agents) and a medium instance for the MySQL service (agent). OpenNebula’s small instances provision VMs equipped with 1 CPU of 2Ghz, and 1GiB of memory, while medium instances are equipped with 4 CPU’s of 2Ghz, and 4GB of memory.

Figure 3 and Figure 4 depict the degree of SLA fulfillment for the naive and history-aware algorithms, indicating the average of response times obtained




Figure 3: PHP resp. time on DAS4 – Naive.

during the execution of the Wikipedia workload trace. The results show how the naive provisioning algorithm tends to generate more SLA violations due to its excessive reactive behavior. These violations are comprised between 700ms (see the yellow Line) and 1500ms (see the red Line) response time values. As we mentioned, this algorithm is an easy target to flash crowds effects, as new VMs can be added or removed to handle sharp and sudden variations in the workload. In contrast on Figure 4, the system performance (*i.e.*, response time) do not fluctuate greatly showing a more stable behavior during the whole experiment. As well we also appreciate a reduction in the number of SLA violations regarding the total amount of PHP-served pages.

Nevertheless to better understand the behavior of both algorithms, we may focus on the resource consumption, as depicted on Figure 5. Firstly, the excessive reactive behavior of the naive algorithm is again illustrated at the interval $t=350min$ and $t=820min$, where two scaling operations under-provision the system during a short period of time. These provisioning decisions provoked fluctuations in the system performance that incremented the financial cost, and therefore throughput alterations. When using the history-aware algorithm, the system makes provisioning decisions by analyzing workload’s trend during

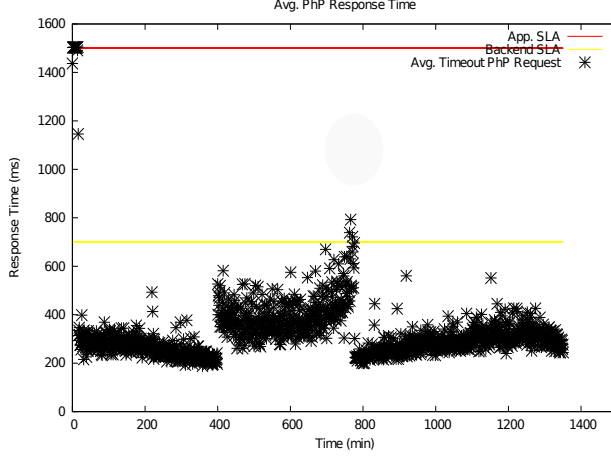


Figure 4: PhP resp. time on DAS4 – History-aware.

a considerable interval of time. Scaling actions are only triggered when having constant alterations in the Wikipedia workload, providing a more efficient resource usage. Indeed, the workload alterations depicted on Figure 1, match with the provisioning decisions made by the history-aware algorithm on Figure 5.

FIXME: I will run more time the History-aware exp. on DAS4.

5.2.1 Discussion

Using the naive provisioning algorithm, the system performance fluctuates greatly following a pattern similar to the web traffic that increases the number of SLA violations. The reactive behavior of this algorithm triggers scaling actions that affect to the system performance instead of improving it, and as a consequence they are also wasteful in terms of resource consumption. Unlike history-aware algorithm offers an efficient resource usage and a constant performance behavior while meeting the application’s SLA. Therefore this algorithm finds the tradeoff between accuracy and cost savings.

Both algorithms are best-effort regarding the SLA fulfillment, and thereby temporal alterations of the workload (with a duration of 5min approx.) cannot




Figure 5: Resource consumption on DAS4.

be guaranteed. The heterogeneity of the PhP-served pages, containing images and requiring multiple Db queries, and the startup time of VMs are in part responsible of these SLA violations.

5.3 Private Cloud

Our experiments on EC2 used small instances for the PhP service (manager and agents) and a medium instance for the MySQL service (agent). EC2 small instances provision VMs equipped with 1 EC2 CPU, and 1.7GiB of memory, while medium instances are equipped with 2 EC2 CPU’s, and 3.75GiB of memory.

In the following, we analysis the behavior of our algorithms when making provisioning decisions on a heterogeneous infrastructure. Figure 6, Figure 7 and Figure 8 show the system performance of the naive, history-aware and profiling-based algorithms, respectively. As depicted on Figure 6, the performance fluctuates greatly following an irregular pattern when using the naive algorithm. More precisely, two of the three workload peaks caused at $t=300min$ and $t=820min$, are explained from the variations on the Wikipedia workload described on Figure 1. However, there is a third peak that corresponds to the interval of time on which the workload trace significantly decreases the number of user’s requests. This explains the degradation of the SLA fulfillment as an effect associated to very frequent scaling actions.




Figure 6: PhP response time on EC2 – Naive.

On the other hand, Figure 7 and Figure 8 show as the history-aware and profiling-based algorithm behave similarly. Even though both algorithms are best-effort, there is an important reduction in the number of SLA violations during the whole experiment. Like on DAS-4, the history-aware algorithm follows a constant performance pattern without having sharp and sudden alterations. Besides, as shown on Figure 8, the profiling-based algorithm has a similar behavior to the history-aware algorithm in terms of system performance, however. The profiling-based algorithm improves the SLA fulfillment in a 11% in comparison with the history-aware algorithm. Therefore we demonstrate how the use of online profiling techniques although intrusive do not cause time delays or throughput alterations.

FIXME: performance or responsiveness

The resource usage on EC2 presents important alterations as shown on Figure 9. When using the naive provisioning, the fluctuations in the system performance are explained as a result of a high frequency of scaling operations. In concrete, these fluctuations caused at the interval of time comprised between $t=400min$ and $t=500min$ (see on Figure 6) match with the provisioning decisions made during the same interval of time on Figure 9. If we now pay attention




Figure 7: PhP resp. time on EC2– History-aware.

to the history-aware, and profiling-based algorithms, their resource consumptions are identical during the whole experiment. Indeed, both algorithms decided to scale out the system during the interval of time comprised between $t=1050min$ and $t=1400min$, to prevent future SLA violations. It demonstrates the benefits of using flexible threshold ranges to provide a predictive provisioning mechanism, improving the user experience.

5.3.1 Discussion

FIXME: Do you prefer this organization for the discussions ? One inside of each type of cloud, and one to summarize.

5.4 Discussion

Generally, the result of our measurements show how the behavioral performance pattern and the resource consumption vary depending on the infrastructures on which we ran our experiments. Different hardware configurations such as those provided by DAS-4 and EC2, offer two distinct scenarios to validate our provisioning algorithms. In these experiments, we demonstrate how trigger-based provisioning mechanisms can affect the system performance instead of




Figure 8: PHP resp. time on EC2– Profiling-based.

improving it, as well as are wasteful in terms of resource usage. Online profiling techniques, although intrusive, were used without producing performance alterations, in fact they slightly reduced the number of SLA violations in comparison with the history-aware algorithm. We also show the benefits by using history-aware and profiling-based provisioning algorithms which find the tradeoff between the accuracy and cost savings.

However, in these experiments, the flexible threshold ranges were pre-defined before execution for all VMs. These threshold values have to change depending the type of instance to be provisioned. Therefore, we believe that offline profiling techniques may be used to define these values depending of the type of instance, thus improving the effectiveness of our predictions.

6 Related studies

There is a wide literature on issues related to dynamic resource provisioning for cloud web applications. Different approaches present solutions based on queuing models [16], feedback loops techniques [6], mathematical models [11] or even approaches using neural networks techniques [7]. However, most of these mod-




Figure 9: Resource consumption on EC2.

els require a deep understanding in mathematics or machine learning techniques which are not easily interpreted by non specialists. Besides the traffic in web applications is shaped by a combination of different factors such as diurnal and seasonal cycles, sociological and psychological, that follows an irregular pattern. It makes extremely challenging the design and development of realistic and accurate dynamic provisioning mechanisms.

These well-known drawbacks force to IaaS like Amazon EC2 and Windows Azure [10], or PaaS like RightScale [14] and OpenShift [12], to design simple threshold rule-based auto-scaling systems, instead of relying on approaches from academic research. Unfortunately, these scaling systems are naive, wasteful in terms of resource consumption and cost savings, and an easy target for flash crowds.

In the following, we present some of the most relevant and realistic academic approaches that proposed dynamic resource provisioning mechanisms for multi-tier applications.

In [15, 16], the authors designed and implemented a predictive and reactive provisioning mechanism. They used a queuing model G/G/1 to decide the server pool size to be provisioned, and an admission control mechanism to face extreme workload variations. Offline profiling techniques were employed to gather information about the resource requirements of the incoming requests for each tier, and thereby to

selectively admit/reject requests for the lightweight files. An evaluation using real-traces on a homogeneous infrastructures shows the benefits of this approach when handling flash crowds. Unfortunately, its admission control mechanism incurs into sporadic SLA violations (if the server utilization exceed a pre-defined threshold) reducing the QoS of the service, and therefore affecting user experience. Similarly to the previous work, [17] extended queuing models and transaction mix models to design a predictive and reactive provisioning system. To model the application performance, they integrated proactive control and feedback control methods that dynamically adjusted the CPU capacity allocated to servers. This work only considered SLA constraints at the system-level, while others constraints at application-specific level such as response time and request rate were not taken into consideration. Besides, an evaluation of CPU variations on a homogeneous infrastructure, when processing traces from a non real-world application, lack arguments to valid its approach.

Regarding the management of flash crowds [19], a proactive application workload manager was designed to separate the user requests between two groups of servers: one named 'base workload' referred to the smaller and smoother variations in the workload; and the other 'trespassing' referred to the temporal burstly workloads caused by flash crowds. To do this, the authors attempt to divide the data items into popular and less popular, and place them in the right group of servers. Even tough a realistic evaluation was conducted on Amazon EC2 utilizing real traces (Yahoo video streaming), authors do not explain in details how the dynamic resource provisioning is done. Recently, in [9], online profiling techniques have been utilized for managing the tradeoff between performance overload, and cost savings for dynamic resource provisioning. The authors replicate at runtime a regular server hosting an application, with a new server with profiling instrumentation. Their experimental results show how profiling techniques can be included in a resource provisioning system, without causing important response time delays or throughput alterations in comparison with non-profiling provisioning. As we mentioned in Sec-

tion 5, profiling techniques can report more benefits than performance degradations or expenses.

7 Conclusion

Excessive reactive algorithms trend to temporally overprovision applications affected by flashcrowds or slashdot effects. This type of algorithms increase the resource consumption and infrastructure costs than other ...

Experiments based on real-traces "Wikipedia" and conducted on heterogeneous and homogeneous cloud infrastructures.

The use of offline-profiling techniques allow to identify the threshold of the resources of a cloud infrastructure.

The provisioning system remains independent of the infrastructure on which the apps run.

8 Acknowledgments

This work is partially funded by the FP7 Programme of the European Commission in the context of the Contrail project under Grant Agreement FP7-ICT-257438.

FIXME: Probably also mention ERRIC here. ??

References

- [1] Advanced School for Computing and Imaging (ASCI). The Distributed ASCI SuperComputer 4. <http://www.cs.vu.nl/das4/>.
- [2] Amazon Elastic Compute Cloud (EC2). <http://aws.amazon.com/ec2>.
- [3] Jiang Dejun, Guillaume Pierre, and Chi-Hung Chi. EC2 performance analysis for resource provisioning of service-oriented applications. In *Proceedings of the 3rd Workshop on Non-Functional Properties and SLA Management in Service-Oriented Computing*, November 2009.

- [4] Guido Urdaneta Erik-Jan van Baaren and Guillaume Pierre. Wikibench: A distributed, wikipedia based web application benchmark. 2009.
- [5] Ganglia monitoring system. <http://ganglia.sourceforge.net/>.
- [6] Jiayu Gong and Cheng-Zhong Xu. A gray-box feedback control approach for system-level peak power management. In *2010 39th International Conference on Parallel Processing (ICPP)*, pages 555–564, September 2010.
- [7] Sadeka Islam, Jacky Keung, Kevin Lee, and Anna Liu. Empirical prediction models for adaptive resource provisioning in the cloud. *Future Gener. Comput. Syst.*, 28(1):155–162, January 2012.
- [8] Dejun Jiang. *Performance Guarantees For Web Applications*. Phd thesis, VU University Amsterdam, March 2012.
- [9] Nima Kaviani, Eric Wohlstadter, and Rodger Lea. Profiling-as-a-service: adaptive scalable resource profiling for the cloud in the cloud. In *Proceedings of the 9th international conference on Service-Oriented Computing, IC-SOC’11*, pages 157–171, Berlin, Heidelberg, 2011. Springer-Verlag.
- [10] Microsoft Windows- Windows Azure. <http://www.windowsazure.com/en-us>.
- [11] Sireesha Muppala, Xiaobo Zhou, Liqiang Zhang, and Guihai Chen. Regression-based resource provisioning for session slowdown guarantee in multi-tier internet servers. *Journal of Parallel and Distributed Computing*, 72(3):362–375, March 2012.
- [12] OpenShift. <https://openshift.redhat.com/app/flex>.
- [13] Guillaume Pierre and Corina Stratan. Con-PaaS: a platform for hosting elastic cloud applications. *IEEE Internet Computing*, 16(5):88–92, September-October 2012.
- [14] RightScale. <http://www.rightscale.com/>.
- [15] Bhuvan Urgaonkar and Prashant Shenoy. Cat-actylsm: Scalable overload policing for internet applications. *Journal of Network and Computer Applications*, 31(4):891–920, November 2008.
- [16] Bhuvan Urgaonkar, Prashant Shenoy, Abhishek Chandra, Pawan Goyal, and Timothy Wood. Agile dynamic provisioning of multi-tier internet applications. *ACM Trans. Auton. Adapt. Syst.*, 3(1):1–39, March 2008.
- [17] Zhikui Wang, Yuan Chen, D. Gmach, S. Singhal, B.J. Watson, W. Rivera, Xiaoyun Zhu, and C.D. Hyser. AppRAISE: application-level performance management in virtualized server environments. *IEEE Transactions on Network and Service Management*, 6(4):240–254, December 2009.
- [18] Wikipedia Foundation. MediaWiki a free software wiki package. <http://www.mediawiki.org>.
- [19] Hui Zhang, Guofei Jiang, Kenji Yoshihira, Haifeng Chen, and Akhilesh Saxena. Resilient workload manager: taming bursty workload of scaling internet applications. In *Proceedings of the 6th international conference on Autonomic computing, ICAC ’09*, pages 45–46, New York, NY, USA, 2009. ACM.